



Lab 03: Making reports with Pandoc

My Favorite Course

Mateusz Mazur (CWL 1)

December 2, 2025

Contents

Introduction	1
Transactions in PostgreSQL	2
Starting a Transaction	2
Committing a Transaction	2
Rolling Back a Transaction	2
Example: Managing Bank Accounts	2
Step 1: Create the Table	2
Step 2: Insert Sample Data	2
Step 3: Perform a Transaction	2
Step 4: Verify the Results	3
Problems in Transactions	3
Non-repeatable Reads	3
Conclusion	4

[@NOTE]: Please note that this text was AI-generated

Introduction

Your task is to create a report demonstrating sample transactions in PostgreSQL. The report should include code snippets and explanations of the SQL commands used. You can use the provided template to structure your report.

This report provides an overview of sample transactions in PostgreSQL. The goal is to demonstrate the use of SQL commands for managing data in a relational database.

Transactions in PostgreSQL

What are transactions in PostgreSQL, and how do you use them? Provide examples of starting a transaction, committing a transaction, and rolling back a transaction.

Transactions in PostgreSQL ensure that a series of operations are executed in a reliable and consistent manner. They follow the ACID properties: Atomicity, Consistency, Isolation, and Durability.

Starting a Transaction

To begin a transaction, use the `BEGIN` command:

```
BEGIN;
```

Committing a Transaction

To save the changes made during a transaction, use the `COMMIT` command:

```
COMMIT;
```

Rolling Back a Transaction

If an error occurs or you want to discard changes, use the `ROLLBACK` command:

```
ROLLBACK;
```

Example: Managing Bank Accounts

Show the use of transactions in a sample scenario, such as transferring money between two bank accounts. Include SQL code snippets and explanations. Provide a plot of the database schema.

This section demonstrates a sample transaction for transferring money between two bank accounts.

Step 1: Create the Table

First, create a table to store account information:

```
CREATE TABLE accounts (
    account_id SERIAL PRIMARY KEY,
    account_name VARCHAR(50),
    balance NUMERIC(10, 2)
);
```

Final database schema is presented in fig. 1.

Step 2: Insert Sample Data

Insert initial data into the `accounts` table:

```
INSERT INTO accounts (account_name, balance)
VALUES
    ('Alice', 1000.00),
    ('Bob', 500.00);
```

Step 3: Perform a Transaction

Transfer \$200 from Alice's account to Bob's account:



Figure 1: Schema of the database used in the exercise

```
BEGIN;

UPDATE accounts
SET balance = balance - 200
WHERE account_name = 'Alice';

UPDATE accounts
SET balance = balance + 200
WHERE account_name = 'Bob';

COMMIT;
```

Step 4: Verify the Results

Check the updated balances:

```
SELECT * FROM accounts;
```

Problems in Transactions

Discuss common problem(s) that can occur in transactions, such as dirty reads, non-repeatable reads, and phantom reads. Provide examples and SQL code snippets to illustrate these issues.

Non-repeatable Reads

Non-repeatable reads occur when a transaction reads the same row twice and gets different data each time due to another transaction modifying that row in between the reads, e.g.:

```
T1: BEGIN;
T1: SET TRANSACTION ISOLATION LEVEL
T1:     READ COMMITTED; -- (default)
```

```

T1: -- T1 reads the data for the first time.
T1: SELECT * FROM accounts
T1:   WHERE account_name = 'Alice';
T1: -- Result: 'Alice', 1000.00

T1: -- T1 reads the data again (just checking).
T1: SELECT * FROM accounts
T1:   WHERE account_name = 'Alice';
T1: -- Result: 'Alice', 1000.00 (the same)

T1: -- T1 reads the data for the second time
T1: --      within the SAME transaction.
T1: SELECT * FROM accounts
T1:   WHERE account_name = 'Alice';
T1: -- Result: 'Alice', 800.00
T1: -- The value has changed from 1000.00 to 800.00,
T1: -- constituting a "Non-repeatable read".
T1: COMMIT;

T2: BEGIN;

T2: -- T2 updates the data for the same row.
T2: UPDATE accounts SET balance = balance - 200
T2: WHERE account_name = 'Alice';

T2: -- T2 commits the change to the database.
T2: COMMIT;

```

Conclusion

Summarize the key points covered in the report about transactions in PostgreSQL.

- This report demonstrated the use of transactions in PostgreSQL to ensure data consistency during operations.
- Transactions are a powerful feature for managing complex database operations safely and reliably.
- Common problems such as non-repeatable reads can occur, and understanding these issues is crucial for effective database management.
- By following best practices and using transactions appropriately, developers can ensure the integrity of their data in PostgreSQL databases.